

# VeriFast

Software Fiável

Mestrado em Engenharia Informática

Mestrado em Informática

Faculdade de Ciências da Universidade de Lisboa

2019/2020

Vasco T. Vasconcelos

# VeriFast

- ▶ VeriFast is a program verifier for Java
- ▶ Uses the design-by-contract approach to modular verification
- ▶ VeriFast is based on separation logic, an extension of the Hoare logic

# Verification

- ▶ Java methods are annotated with pre and postconditions and other specifications describing assumptions made by the developer
- ▶ VeriFast checks whether the assumptions hold in each execution of the program for arbitrary input
- ▶ If VeriFast deems a Java program to be correct, then that program
  - ▶ does not contain assertion violations
  - ▶ data races
  - ▶ divisions by zero
  - ▶ null dereferences
  - ▶ array indexing errors
  - ▶ and the program makes correct use of the Java API

# Contracts

- ▶ Conventional Java code is not analysed by VeriFast
- ▶ In order to call VeriFast's attention add a contract to a method:

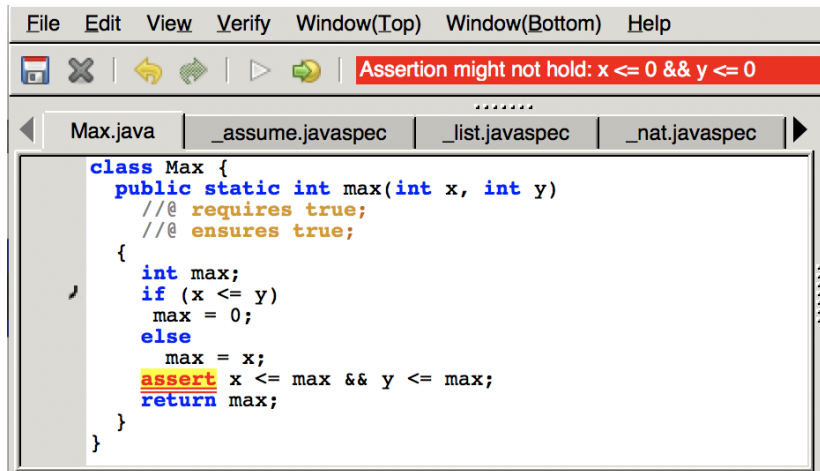
```
static int toInt (Integer i)
    //@ requires true;
    //@ ensures true;
{
    return i.intValue();
}
```

- ▶ A contract for a method is a pair **requires/ensures** placed in line comments, `//@`, or block comments, `/*@ ... @*/`

# The assert statement

- ▶ A Java assert statement consists of the keyword `assert` followed by a boolean expression
- ▶ By inserting an `assert` statement in the code, a developer indicates that she expects the corresponding boolean expression to evaluate to true whenever the statement is reached during the program's execution
- ▶ If the expression evaluates to false, an `AssertionError` is thrown (provided assertion checking is enabled)
- ▶ VeriFast, on the other hand, checks `assert` statements *without* evaluating any code

## Max



The screenshot shows an IDE window with a menu bar (File, Edit, View, Verify, Window(Top), Window(Bottom), Help) and a toolbar with icons for file operations and execution. A red status bar at the top displays the message: "Assertion might not hold:  $x \leq 0 \ \&\& \ y \leq 0$ ". The file explorer shows four files: Max.java, \_assume.javaspec, \_list.javaspec, and \_nat.javaspec. The main editor displays the code for Max.java:

```
class Max {  
    public static int max(int x, int y)  
        //@ requires true;  
        //@ ensures true;  
    {  
        int max;  
        if (x <= y)  
            max = 0;  
        else  
            max = x;  
        assert x <= max && y <= max;  
        return max;  
    }  
}
```

# Method contracts

- ▶ VeriFast performs modular verification: each method call is verified with respect to the callee's signature
- ▶ The current contract of method `max`, namely

```
//@ requires true;  
//@ ensures true;
```

tells very little about the behaviour of the method

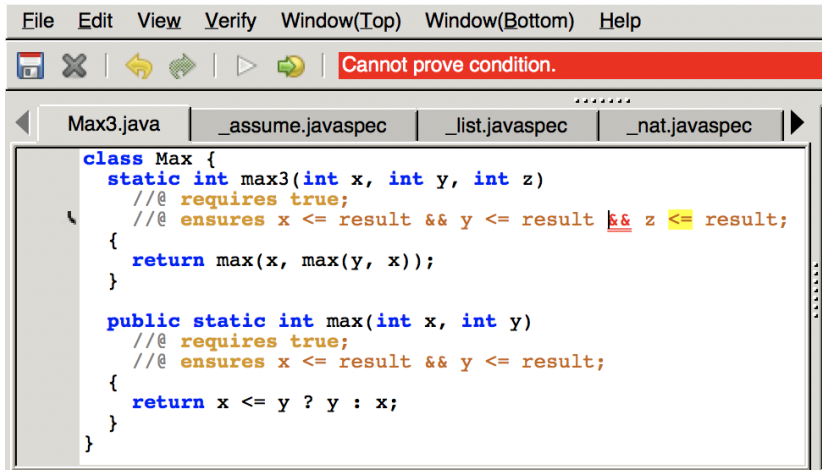
- ▶ Instead we “move” the `assert` expression to the post condition,

```
//@ ensures x <= result && y <= result
```

so that we may use the contract in another method

- ▶ Notice the *ghost variable* `result` used to denote the value of the method

## Max3



The screenshot shows a Java IDE window with a menu bar (File, Edit, View, Verify, Window(Top), Window(Bottom), Help) and a toolbar with icons for saving, undo, redo, and running. A red error bar at the top of the editor area displays the message "Cannot prove condition." The editor has four tabs: Max3.java, \_assume.javaspec, \_list.javaspec, and \_nat.javaspec. The Max3.java tab is active, showing the following code:

```
class Max {  
    static int max3(int x, int y, int z)  
        //@ requires true;  
        //@ ensures x <= result && y <= result && z <= result;  
    {  
        return max(x, max(y, x));  
    }  
  
    public static int max(int x, int y)  
        //@ requires true;  
        //@ ensures x <= result && y <= result;  
    {  
        return x <= y ? y : x;  
    }  
}
```



## Partial correctness

- ▶ A method body *satisfies a contract* if for each program state  $s$  that satisfies the precondition, execution of the method body starting in  $s$  does not trigger illegal operations (such as assertion violations and divisions by zero) and the postcondition holds when the method terminates
- ▶ VeriFast *only checks partial correctness* so methods are not required to terminate

# Symbolic execution

- ▶ VeriFast uses *symbolic* rather than concrete execution
- ▶ It constructs a symbolic state that represents an arbitrary concrete pre-state which satisfies the precondition
- ▶ Checks that the body satisfies the contract for this symbolic state
- ▶ Symbolically executes the body starting in the initial symbolic state
- ▶ At each statement encountered during symbolic execution, checks that the statement cannot go wrong and updates the symbolic state to reflect execution of that statement
- ▶ Finally, when the method returns, VeriFast checks that the postcondition holds for *all* resulting symbolic states

# Symbolic state

- ▶ A symbolic state is a triple composed of
  - ▶ A *symbolic store* (right frame on the IDE)
  - ▶ A path condition, or *assumptions* (bottom-centre frame on the IDE)
  - ▶ A *symbolic heap* (right-centre frame on the IDE)
- ▶ Each *symbolic value* is a first-order term, i.e., a symbol, or a literal number, or an operator (+, -, <, =, ...) or a function applied to first-order terms.
- ▶ The path condition is a set of first-order formulas describing the conditions that hold on the path being verified
- ▶ The symbolic heap is a multi-set of heap chunks (more later)

# Assertions

- ▶ An assertion is a side-effect free, heap-independent Java boolean expression (extensions to be introduced later)
- ▶ *Consuming an assertion*—**ensures**, **assert**—means symbolically evaluating the expression yielding a first-order formula and checking that the formula is derivable from the path condition
- ▶ VeriFast relies on an SMT solver, a kind of automatic theorem prover, to discharge such proof obligations
- ▶ *Producing an assertion*—**requires**, **assume**—corresponds to evaluating that expression yielding a first-order formula and adding it to the path condition

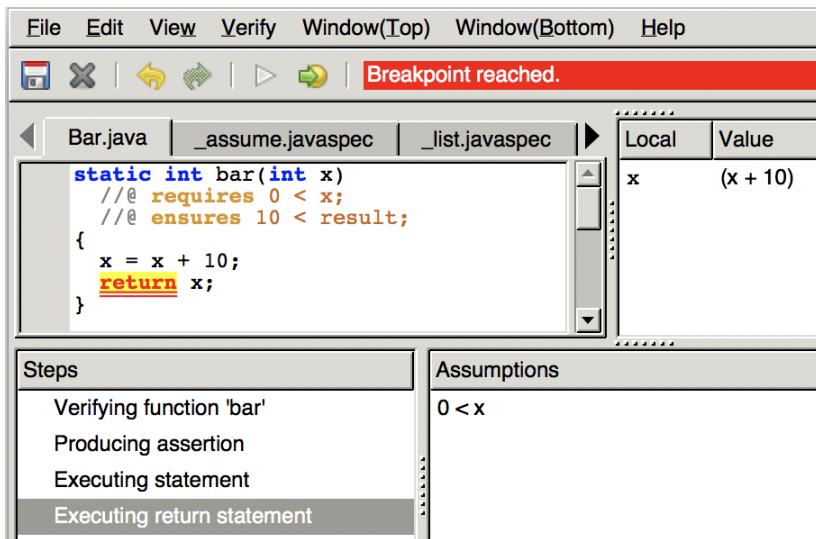
# The pre-state

- ▶ Symbolic execution of each method starts by initializing the symbolic store by assigning a fresh first-order symbol to each parameter
- ▶ VeriFast selects the symbol  $x$  as the fresh term representing the symbolic value of a parameter  $x$
- ▶ The resulting symbolic state thus represents an arbitrary concrete pre-state
- ▶ Notice that the symbol and its symbolic value are rendered in different fonts:  $x$  and  $x$

# Assignment

- ▶ Disable overflow warnings by unchecking Check arithmetic overflow in the Verify menu
- ▶ Place the cursor in `return` statement
- ▶ Press the “Run to cursor” button
- ▶ Check the Symbolic Store frame (top left) and the Assumptions frame (bottom centre)
- ▶ The postcondition holds as the corresponding first-order formula,  $10 < x + 10$ , is derivable from the path condition,  $0 < x$ , by the SMT solver

# Assignment example



The screenshot shows a Java IDE with a breakpoint reached in the 'bar' function of 'Bar.java'. The IDE has a menu bar (File, Edit, View, Verify, Window(Top), Window(Bottom), Help) and a toolbar with icons for saving, closing, undo, redo, step over, and step into. A red banner at the top of the toolbar area says 'Breakpoint reached.'.

The editor shows the following code in 'Bar.java':

```
static int bar(int x)
    //@ requires 0 < x;
    //@ ensures 10 < result;
{
    x = x + 10;
    return x;
}
```

The 'return' statement is highlighted in yellow. The right-hand pane shows the local variable 'x' with the value '(x + 10)'.

Local	Value
x	(x + 10)

The bottom-left pane shows the 'Steps' of the execution:

- Verifying function 'bar'
- Producing assertion
- Executing statement
- Executing return statement

The bottom-right pane shows the 'Assumptions':

- $0 < x$

# Conditional

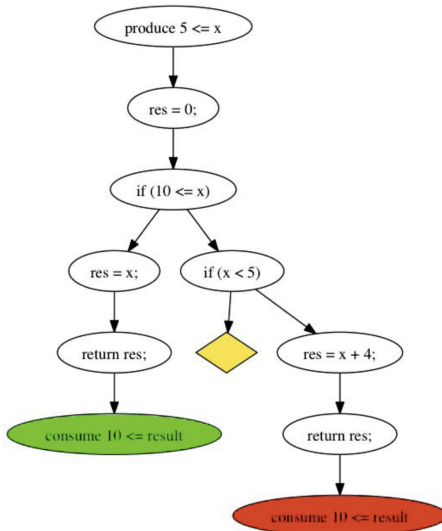
- ▶ Symbolic evaluation of the condition of an if statement results in a first-order formula
- ▶ Based on this formula, it is generally not possible to decide which branch must be taken
- ▶ Given a statement `if (x<10)S1; else S2;`, statement S1 is verified under the assumption that  $10 \leq x$ , while S2 is verified assuming the negation of the condition,  $10 > x$



## Conditional example

```
static int foo(int x)
  //@ requires 5 <= x;
  //@ ensures 10 <= result;
{
  int res = 0;
  if (10 <= x)
    res = x;
  else if (x < 5)
    assert false;
  else
    res = x + 4;
  return res;
}
```

# Symbolic execution tree



## Inconsistent assumptions

- ▶ The diamond node represents a symbolic state with an inconsistent path condition
- ▶ Such states are not reachable during concrete executions of the program
- ▶ VeriFast does not examine infeasible paths any further
- ▶ The formula representing the postcondition,  $10 \leq x + 4$ , is not derivable from the path condition  $5 \leq x$ ,  $10 > x$  and  $x \geq 5$
- ▶ VeriFast does not report all problems on all paths; it stops when it finds the first error or when all paths successfully verify

# Method calls


- ▶ The symbolic execution of a call consists of two steps:
  1. Consumption of the callee's precondition and
  2. Production of its postcondition
- ▶ Both steps are executed under the callee's symbolic store
- ▶ During production of the postcondition, the callee's return value is represented by the ghost variable `result`

# The memory layout of a Java program

- ▶ Before we address classes and objects we need to understand the memory layout of a Java program

```
class MemoryLayout {  
    static String g (boolean b) {  
        return Boolean.toString(!b);  
    }  
    static String f (boolean b) {  
        return g(b).toUpperCase();  
    }  
    public static void main (String[] args) {  
        System.out.println(f(true));  
    }  
}
```

# The memory layout of a Java program



# Aliasing and heap chunks

- ▶ Modular verification in the presence of aliasing is challenging
- ▶ VeriFast applies an *ownership* regime
- ▶ VeriFast tracks during symbolic execution what part of the program state is owned by the method
- ▶ The symbolic heap is a *multiset of heap chunks*
- ▶ Each heap chunk represents a memory region that is owned by a method
- ▶ The chunk can contain information on the *state* of that memory region. For example, the heap chunk  $C.f(o, v)$  represents
  - ▶ exclusive ownership of the field  $C.f$  of object  $o$  and
  - ▶ the property that the field's current value is  $v$

## Heap chunks in a symbolic heap

- ▶ All heap chunks in the symbolic heap represent mutually disjoint memory regions
- ▶ If the heap contains two field chunks  $C_f(o_1, v)$  and  $C_f(o_2, w)$ , then  $o_1$  and  $o_2$  are distinct
- ▶ As chunks on the symbolic heap do not share hidden dependencies, the verifier can safely assume that an operation that only affects a particular chunk does not invalidate the information in the remaining chunks
- ▶ **Invariant:** any time each heap location is exclusively owned by at most one activation record
- ▶ A method is only allowed to access a heap location if it owns that location



# Acquiring and releasing ownership

- ▶ Acquiring
  1. Constructor
  2. Precondition
  3. Call
- ▶ Releasing:
  1. Call
  2. Return

## Aquiring ownership via construction

```
class Account {  
    private int balance;  
    public Account()  
        //@ requires true;  
        //@ ensures balance  $\mapsto$  0;  
    {  
        super();  
        balance = 0;  
    }...}
```

- ▶ A constructor gains ownership of the fields of the new object right after calling the superclass constructor
- ▶ The constructor of the class `Account` is allowed to initialize `balance` to zero

## Aquiring ownership via precondition

```
public void deposit(int amount)
  //@ requires balance  $\mapsto$  ?b;
  //@ ensures balance  $\mapsto$  b + amount;
{
  this.balance += amount;
}
```

- ▶ Ownership of the field `f` of an object `e1` with value `e2` is denoted as `e1.f  $\mapsto$  e2`. Read “`e1.f` points to `e2`”
- ▶ The method body of method `deposit` is allowed to read and write `this.balance`
- ▶ Note the `?b` to introduce a new ghost variable `b`. This is typical of the precondition; the variable can then be used in the postcondition

## Aquiring ownership via call

```
1 public Account copy()
2     //@ requires balance  $\mapsto$  ?b;
3     //@ ensures balance  $\mapsto$  b  $\&*\&$  result != null
4          $\&*\&$  result.balance  $\mapsto$  b;
5 {
6     Account copy = new Account();
7     copy.balance = balance;
8     return copy;
9 }
```

- ▶ A new Account is created in line 5. Assignment to balance (line 6) is allowed because the constructor's postcondition includes `this.balance  $\mapsto$  ...`
- ▶ The postcondition of the constructor specifies that ownership of field balance of the new object is transferred to its caller (copy() in this case) when the constructor terminates

## Releasing ownership via call

```
1 public void release()  
2     //@ requires this.balance  $\mapsto$  _;  
3     //@ ensures true; // balance not released  
4 {}  
5 public void m()  
6     //@ requires this.balance  $\mapsto$  _;  
7     //@ ensures true;  
8 {  
9     release();  
10    this.balance += 0; // Error _ No matching heap  
11                          chunks: Account_balance(this, _)  
12 }
```

- ▶ At each method call, ownership of the memory locations described by the callee's precondition is transferred from the caller to the callee
- ▶ Ownership of `balance` is lost in line 9

## Releasing ownership via return

```
public void deposit(int amount)
    //@ requires balance  $\mapsto$  ?b;
    //@ ensures balance  $\mapsto$  b + amount;
{ this.balance += amount; }
public void m ()
    //@ requires this.balance  $\mapsto$  _;
    //@ ensures true;
{
    deposit(100);
    this.balance += 0; // OK: there is a matching
                       h. chunk released by deposit() in its post
}
```

- ▶ When a method returns, the method loses ownership of all memory locations enumerated in its postcondition
- ▶ Ownership of those locations is transferred from the method to its caller when it returns

# Spatial assertions

- ▶ *Pure assertions* such as  $0 \leq x$  specify constraints on local variables
- ▶ *Spatial assertions* such as  $o.f \mapsto v$  denote ownership of a heap subregion and information about that region
- ▶ *Production of a spatial assertion* corresponds to the acquisition of ownership by the current activation record of the memory regions described by the assertion
- ▶ *Consumption of a spatial assertion* corresponds to the current activation record relinquishing ownership of the memory regions described by the assertion

## Separating conjunction

- ▶ Multiple atomic assertions can be conjoined via the separating conjunction, denoted  $\&*\&$
- ▶ Semantically,  $A \&*\& B$  holds if both  $A$  and  $B$  hold and  $A$ 's footprint is disjoint from  $B$ 's footprint
- ▶ The footprint of an assertion is the set of memory locations for which that assertion claims ownership
- ▶ Consuming (respectively producing)  $A \&*\& B$  is implemented by first consuming (respectively producing)  $A$  and afterwards  $B$
- ▶ Note that if  $A$  is a pure assertion, then  $A \&*\& A$  is equivalent to  $A$ . However, this property does not necessarily hold for spatial assertions



# Ownership transfer

File Edit View Verify Window(Top) Window(Bottom) Help

Breakpoint reached.

Account.java \_assume.javaspec \_list.javaspec \_nat.javaspec

```

public void transfer(Account other, int amount)
/*@ requires this.balance |-> ?b1 &*& //other != null &*&
    other.balance |-> ?b2; @*/
/*@ ensures this.balance |-> b1 - amount &*&
    other.balance |-> b2 + amount; @*/
{
    balance -= amount;
    other.deposit(amount);
}

```

Local	Value
amount	amount
b1	b1
b2	b2
other	other
this	this

Steps

- Verifying method 'transfer'
- Producing assertion
- Producing assertion

Assumptions

- !(this = 0)
- !(other = 0)
- !(other = this)

Heap chunks

- Account\_balance(other, b2)
- Account\_balance(this, b1)
- Account\_balance(this, (b1 - amount))

## Ownership transfer

- ▶ Place the cursor in the method's closing brace; press "Run to cursor"
- ▶ Click on a step (lower-left box) and observe the "Heap chunks" frame; use the up/down key to see the verification progress

## One heap chunk per field

- ▶ Different references may hold different parts of an object
- ▶ Consider the following excerpt of class FieldSeparation

```
int a;  
boolean b;  
int getA ()  
    //@ requires a  $\mapsto$  _;  
    //@ ensures true; // do not release a  
{ return a; }  
boolean getB ()  
    //@ requires b  $\mapsto$  _;  
    //@ ensures true; // do not release b  
{ return b; }
```

## Different references may hold different parts of an object

```
FieldSeparation f = new FieldSeparation();  
f.getA(); f.getB(); // OK
```

```
FieldSeparation f1 = new FieldSeparation();  
FieldSeparation f2 = f1;  
f1.getA(); f2.getB(); // OK
```

```
FieldSeparation f3 = new FieldSeparation();  
f3.getA(); f3.getA(); // No mathing heap chunks:  
    FieldSeparation_a(f3, _)
```

```
FieldSeparation f4 = new FieldSeparation();  
FieldSeparation f5 = f4;  
f4.getA(); f5.getA(); // No mathing heap chunks:  
    FieldSeparation_a(f4, _)
```

# Data abstraction

- ▶ The contract for class Account is built around the value of field balance
- ▶ But balance is a **private** field, hence should not appear in contracts
- ▶ If we choose a different representation for the class (e.g. a list of transactions), we would have to
  - ▶ Update the method contracts and consequently
  - ▶ Have to reconsider the correctness of all clients

# Predicates

- ▶ How can we specify the observable behaviour of a class or interface without exposing its internal representation?
- ▶ VeriFast's answer to this question is **predicates**
- ▶ Assertions describing the state associated with instances of a class can be hidden inside predicates
- ▶ A predicate is a named, parameterised, assertion

```
//@ predicate account(int b) = this.balance  $\mapsto$  b  
    &*& b >= 0;
```

- ▶ We now use predicate account in all method contracts
- ▶ This predicate is defined **within** class Account

## Class Account with predicate account()

```
class Account {  
    private int balance;  
    //@ predicate account(int b) = this.balance ==  
        b && b >= 0;  
    public Account()  
        //@ requires true;  
        //@ ensures account(0);  
    {}  
    public void deposit(int amount)  
        //@ requires account(?b);  
        //@ ensures account(b + amount);  
    { this.balance += amount; }  
}
```

- Contracts are now written in terms of predicate `account()` and not field `balance`

## Static predicates

- ▶ Predicates that talk about possibly **null** references cannot be declared inside a class (**this** != **null**)
- ▶ In this case we declare the predicate **outside** the class and pass the object as parameter

```
//@ predicate account(Account a, int b) = a.  
    balance ↦ b &*& b >= 0;
```

- ▶ Calls to this predicate require a new parameter, typically **this**:

```
public void deposit(int amount)  
    //@ requires account(this, ?b);  
    //@ ensures account(this, b + amount);
```



## Predicates are class members

- ▶ A call to a member `account()` abbreviates `this.account()`
- ▶ To call the other predicate on an object `o` use `o.account()`

```
public void transfer(Account target, int amount)
    /*@ requires amount >= 0 &*&
        this.account(?b1) &*& amount <= b1 &*&
        target != null &*& target.account(?b2);
        @*/
    /*@ ensures this.account(b1 - amount) &*&
        target.account(b2 + amount); @*/
```

# Folding and unfolding predicates

- ▶ VeriFast by default does not automatically fold and unfold predicates
- ▶ Developers must explicitly use ghost statements to switch between the external, abstract view offered by the predicate and the internal definition of the predicate

## Folding (closing) a predicate

- ▶ The `close` ghost statement *folds* a predicate: it consumes the body of the predicate, and afterwards adds a chunk representing the predicate to the symbolic heap (see method `Account`)
- ▶ Without the ghost statement, the constructor does not verify as the heap does not contain a chunk that matches the postcondition

```
public Account(int initialBalance)
    //@ requires initialBalance >= 0;
    //@ ensures account(this, initialBalance);
{
    balance = initialBalance;
    //@ close account(this, initialBalance);
}
```

## Unfolding (opening) a predicate

- ▶ The **open** ghost statement unfolds a predicate: it removes a heap chunk that represents the predicate from the symbolic heap and produces its body
- ▶ As the necessary chunk is nested inside `account(this, ?b)`, the predicate must be opened first
- ▶ If we omit the ghost statement, VeriFast would no longer find a chunk that matches the field assertion `Account_balance(this, _)` on the heap and report an error

```
public void deposit(int amount)
  //@ requires account(this, ?b);
  //@ ensures account(this, b + amount);
{
  //@ open account(this, b);
  this.balance += amount;
  //@ close account(this, b + amount);
}
```

## Precise predicates

- ▶ Inserting **open** and **close** ghost statements is tedious
- ▶ To alleviate this burden, programmers can mark certain predicates as **precise**
- ▶ VeriFast automatically opens and closes precise predicates (in many cases) whenever necessary during symbolic execution
- ▶ A predicate can be marked as precise by using a semicolon instead of comma somewhere in the parameter list

```
//@ predicate account(Account a; int b) = a.  
    balance  $\mapsto$  b;
```

- ▶ The semicolon separates the input from the output parameters: a is input; b is output

## Account with precise predicates

- ▶ No **open** or **close** ghost statements required
- ▶ Nevertheless, if possible declare predicates inside the class

```
class Account {  
    private int balance;  
    public Account()  
        //@ requires true;  
        //@ ensures account(this, 0);  
    {}  
    public void deposit(int amount)  
        //@ requires account(this, ?b);  
        //@ ensures account(this, b + amount);  
    { this.balance += amount; }  
}
```

# An alternative implementation for class Account

- ▶ Now that we have the contract for class `Account` defined on top of a predicate, we can easily change the implementation without touching the contract
- ▶ The implementation:
  - ▶ Stores the deposit/withdraw transactions on a linked list,
  - ▶ Extracts the balance from the list of transactions, rather than storing it explicitly on a field, and
  - ▶ Introduces a new definition for `predicate` `account`

# Class Transaction

```
class Transaction {  
    final Transaction next;  
    final int amount;  
    public Transaction(int a, Transaction t)  
        //@ requires true;  
        //@ ensures amount  $\mapsto$  a &* & next  $\mapsto$  t;  
    { amount = a; next = t; }  
}
```

- Class Transaction implements a linked list of integer values



## Account with a list of transactions

```
class Account {  
    private Transaction transactions;  
    public Account()  
        //@ requires true;  
        //@ ensures account(0);  
    {}  
    ...  
}
```

- ▶ The contract for the constructor (and other methods) remains unchanged

## The new predicate for the class

```
/*@  
predicate account(int b) =  
    this.transactions  $\mapsto$  ?ts &* &  
    transactions(ts, b);  
@*/
```

- ▶ Field transactions points to ts and predicate transactions (below) is true of ts and balance b
- ▶ Again, note ?ts to introduce variable ts

## The predicate for class Transaction

```
/*@  
predicate transactions(Transaction t; int total)  
=  
t == null ?  
  b == 0  
:  
  t.amount  $\mapsto$  ?a &* &  
  t.next  $\mapsto$  ?n &* &  
  transactions(n, ?ntotal) &* &  
  total == a + ntotal;  
@*/
```

- ▶ transactions() is a recursive predicate that traverses the list collecting the amounts in the transactions
- ▶ Declared outside the class so that it may be used with `null` references

## The balance of an Account with transactions

```
public int getBalance ()
    //@ requires account(?b);
    //@ ensures account(b);
{
    return getTotal(transactions);
}

private int getTotal(Transaction t)
    //@ requires transactions(t, ?total);
    //@ ensures transactions(t, total) &*& result
        == total;
{
    //@ open transactions(t, total);
    return t == null ? 0 : t.amount + getTotal(t.
        next);
}
```

- In this case an **open** transactions() is mandatory for the success of Verifast

# Inheritance

- ▶ A Java interface defines a set of methods; each non-abstract class that implements the interface provides code for each method
- ▶ In order to modularly verify client code, each interface method is annotated with a contract
- ▶ We have used predicates to hide the internal representation of classes
- ▶ Verifast interfaces allows **predicate** declarations in interfaces
- ▶ A class that implements an interface is a subtype of the interface; contracts for subtypes can be refined
- ▶ In any case they must be restated in the subtype (even if they remain unchanged)

## A stack interface that only speaks about the size

```
interface StackSize {
    //@ predicate stack(int size);
    void push(Object o);
        //@ requires stack(?s);
        //@ ensures stack(s+1);
    int size();
        //@ requires stack(?s);
        //@ ensures stack(s) &* & result == s;
    Object peek ();
        //@ requires stack(?s) &* & s > 0;
        //@ ensures stack(s);
    void pop ();
        //@ requires stack(?s) &* & s > 0;
        //@ ensures stack(s - 1);
}
```

# Array ownership

- ▶ Recall: ownership of field  $f$  of an object  $e1$  with value  $e2$  is denoted as  $e1.f \mapsto e2$ . Read “ $e1.f$  points to  $e2$ ”
- ▶ For arrays one writes:  $a[\text{from}..\text{to}] \mapsto v$  to mean “the portion of array  $a$  between indices  $\text{from}$  (inclusive) to  $\text{to}$  (exclusive) points to  $v$ ”
- ▶ Example for a stack implemented with an array `elements`, variable `elems` denotes the list of the elements in the stack:

```
this.size  $\mapsto$  ?s &* & this.elements  $\mapsto$  ?e &* &  
e[0..s]  $\mapsto$  elems &* & ...
```

# ArrayStack

```
public class ArrayStack implements StackSize {  
    private Object[] elements;  
    private int size;  
    /*@  
    predicate stack(int s) =  
        this.size  $\mapsto$  s &*&           // Acquire size  
        this.elements  $\mapsto$  ?e &*&    // Acquire reference  
        e[0..e.length]  $\mapsto$  _ &*&    // Acquire all array  
        elems  
        0 <= s && s <= e.length; // The invariant  
    @*/
```



## Some ArrayStack methods

```
ArrayStack(int initialCapacity)
    //@ requires initialCapacity >= 0;
    //@ ensures stack(0);
{ elements = new Object[initialCapacity]; }

void pop()
    //@ requires stack(?s) &* & s > 0;
    //@ ensures stack(s - 1);
{ elements[--size] = null; }

Object peek()
    //@ requires stack(?s) &* & s > 0;
    //@ ensures stack(s);
{ return elements[size - 1]; }
```

- Contracts written with **predicate** stack() only

## Copying an array

```
Object[] a = new Object [size * 2 + 1];
//@ close array_slice_dynamic(array_slice_Object
    , elements, 0, size, _); // get hold of the
    elems array
//@ close array_slice_dynamic(array_slice_Object
    , a, 0, size, _); // same for array a
//@ close arraycopy_pre(array_slice_Object,
    false, 1, elements,
0,size, _, a, 0); // fold the pre
System.arraycopy(elements, 0, a, 0, size);
//@ open arraycopy_post(_, _, _, _, _, _, _, _
    _); // unfold the post
//@ open array_slice_dynamic(array_slice_Object,
    a, _, _, _); // release array a
elements = a;
```

## Method push: growing the stack

```
void push(Object x)
    //@ requires stack(?s);
    //@ ensures stack(s + 1);
{
    if (size == elements.length) {
        Object[] a = new Object [size * 2 + 1];
        ...
        System.arraycopy(elements, 0, a, 0, size);
        ...
    }
    elements[size++] = x;
}
```

## A better invariant

- ▶ The invariant for the stack, and consequently the contracts, only talk about the size
- ▶ We never know if the `push()` indeed places the value in the stack, let alone if the value is placed at the top
- ▶ For more precise invariants, we use **models**

# Model-based specifications

- ▶ Modelling is an abstraction technique for system design and specification
- ▶ A **model** is a representation of the desired system
- ▶ A **formal model** is one that has a precise description in a formal language
- ▶ A model differs from an implementation in that it might:
  - ▶ capture only some aspects of the system
  - ▶ be partial, leaving some parts unspecified
  - ▶ not be executable
- ▶ An implementation of the system can be compared to the model

## The list datatype: one of the Verifast models

- ▶ See <verifast>/bin/rt/\_list.javaspec
- ▶ Operations on datatypes (introduced with the keyword `fixpoint`) must be total

```
inductive list<t> = nil | cons(t, list<t>);
```

```
fixpoint t head<t>(list<t> xs) {  
  switch (xs) {  
    case nil: return default_value<t>;  
    case cons(x, xs0): return x;  
  }  
}
```

## Some list predefined operations

```
fixpoint t head<t>(list<t> xs)
fixpoint list<t> tail<t>(list<t> xs)
fixpoint int length<t>(list<t> xs)
fixpoint list<t> append<t>(list<t> xs, list<t> ys)
fixpoint list<t> reverse<t>(list<t> xs)
fixpoint boolean mem<t>(t x, list<t> xs)
fixpoint t nth<t>(int n, list<t> xs)
fixpoint list<t> store<t>(list<t> xs, int index, t v)
fixpoint boolean distinct<t>(list<t> xs)
fixpoint list<t> take<t>(int n, list<t> xs)
fixpoint list<t> drop<t>(int n, list<t> xs)
fixpoint list<t> remove<t>(t x, list<t> xs)
fixpoint list<t> remove_nth<t>(int n, list<t> xs)
fixpoint list<t> remove_every<t>(t x, list<t> xs)
fixpoint list<t> remove_all<t>(list<t> xs, list<t> ys)
fixpoint int index_of<t>(t x, list<t> xs)
fixpoint boolean all_eq<t>(list<t> xs, t x0)
fixpoint list<t> update<t>(int i, t y, list<t> xs)
fixpoint boolean forall<t>(list<t> xs, fixpoint(t,
    boolean) p)
```

## An interface that talks about the elements in the stack

```
interface Stack {  
    //@ predicate stack(list<Object> elems);  
    void push(Object x);  
        //@ requires stack(?e);  
        //@ ensures stack(append(e, cons(x, nil)));  
    int size();  
        //@ requires stack(?e);  
        //@ ensures stack(e) &*& result == length(e)  
        ;  
    Object peek ();  
        //@ requires stack(?e) &*& length(e) > 0;  
        //@ ensures stack(e) &*& result == nth(  
            length(e) - 1, e);  
    void pop ();  
        //@ requires stack(?e) &*& length(e) > 0;  
        //@ ensures stack(take(length(e) - 1, e));  
}
```



# The abstraction function for an ArrayStack

```
public class ArrayStack implements Stack {
    /*@
    predicate stack(stack<Object> elems) =
        this.size  $\mapsto$  ?s &*&
        this.elements  $\mapsto$  ?e &*&
        e[0..s]  $\mapsto$  elems &*& // get hold of the
            elems in the stack
        e[s..e.length]  $\mapsto$  _ &*& // get hold of the
            remaining elems in the array
        s == length(elems); // the invariant
    @*/
```

## Some ArrayStack methods

```
ArrayStack(int initialCapacity)
//@ requires initialCapacity >= 0;
//@ ensures stack(nil);

boolean isEmpty ()
//@ requires stack(?elems);
//@ ensures stack(elems) &*& result == (length(
    elems) == 0);
```

# The push method

```
void push(Object x)
  //@ requires stack(?elems);
  //@ ensures stack(append(elems, cons(x, nil)));
  // Cannot prove condition.
```

► Recall the invariant:

```
this.size  $\mapsto$  ?s &*&
this.elements  $\mapsto$  ?e &*&
e[0..s]  $\mapsto$  elems &*& // get hold of the stack
    elems
s == length(elems); // the invariant
```

## The verification condition for push()

- ▶ In this case, `elems` in the invariant is a complicated expression

```
append(elemsOld, cons(x, nil))
```

where `elemsOld` is the list at method entry

- ▶ So Verifast needs to prove that

```
s == length(append(elemsOld, cons(x, nil)));
```

- ▶ that is, that the length of the append of two lists is the sum of the length of the lists:

```
length(append(xs, ys)) == length(xs) +  
    length(ys)
```

- ▶ And this is a bit too much for VeriFast

# Helping Verifast

- ▶ Lemma functions allow developers to prove properties about predicates
- ▶ A lemma is a method without side effects. A lemma is *pure* if its contract does not contain spatial assertions
- ▶ A particular useful variant is the `lemma_auto` that does not require to write a body:

```
/*@  
lemma_auto void length_append<t>(list<t> xs,  
    list<t> ys)  
    requires true;  
    ensures length(append(xs, ys)) == length(xs) +  
        length(ys);  
{  
    length_append(xs, ys);  
}  
@*/
```

# Inheritance

- ▶ Contracts are not inherited: there exists a relation between contracts in the subtype and that of the supertype
- ▶ In particular, the two contracts may coincide

```
interface Parent {  
    int triple(int n);  
    //@ requires n >= 0;  
    //@ ensures result >= 0;  
}  
  
class Child1 implements Parent {  
    int triple(int n)  
        //@ requires n >= 0;           // As in supertype  
        //@ ensures result >= 0;     // As in supertype  
    { return 3 * n; }  
}
```

## The precondition may be weakened

```
class Child2 implements Parent {  
    int triple(int n)  
        //@ requires n > -5;  
        //@ ensures result >= 0;  
    { return n > 0 ? 3 * n : -3 * n; }  
}
```

- Precondition  $n > -5$  **implies**  $n > 0$ , and so we are good

## The postcondition may be strengthened

```
class Child3 implements Parent {  
    int triple(int n)  
        //@ requires n >= 0;  
        //@ ensures result >= n;  
    { return 3 * n; }  
}
```

- Postcondition `result >= n` is implied by `result >= 0`, and so we are good



## Weak the pre; strengthen the post

```
class Child4 implements Parent {  
    int triple(int n)  
        //@ requires true;  
        //@ ensures result >= n;  
    { return n > 0 ? 3 * n : -3 * n; }  
}
```

- ▶ The precondition in the subtype implies that in supertype
- ▶ The postcondition in the subtype is implied by that in supertype

# Bibliography



Jan Smans, Bart Jacobs, and Frank Piessens.

Verifast for Java: A tutorial.

In *Aliasing in Object-Oriented Programming. Types, Analysis and Verification*, volume 7850 of *Lecture Notes in Computer Science*, pages 407–442. Springer, 2013.